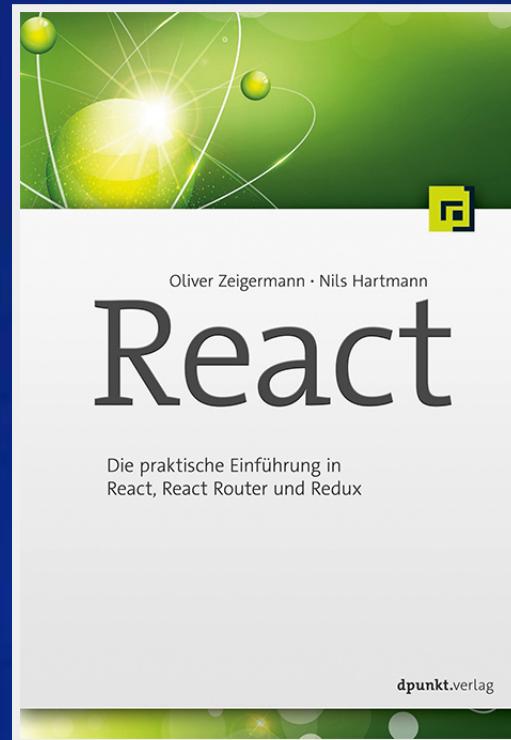
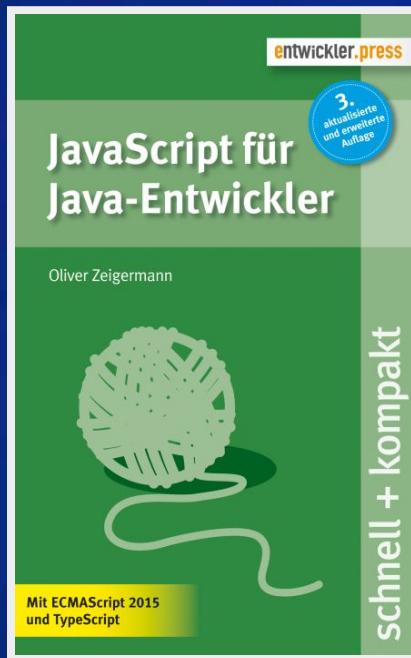


Machine Learning in the Browser with Deep Neural Networks

data2day 2016

Oliver Zeigermann / @DJCordhose

<http://bit.ly/ml-data2day>



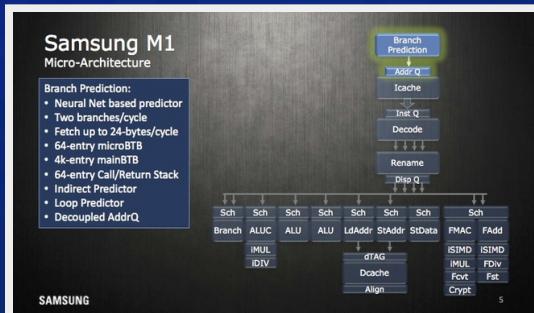
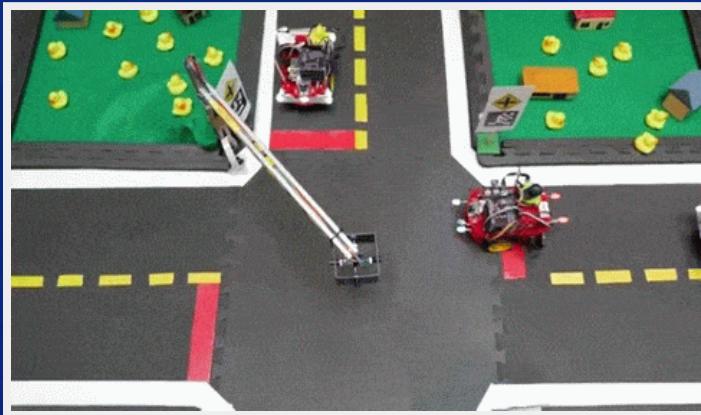
embarc



What is machine learning?

**science of getting computers to act
without being explicitly programmed**

Applications of machine learning?



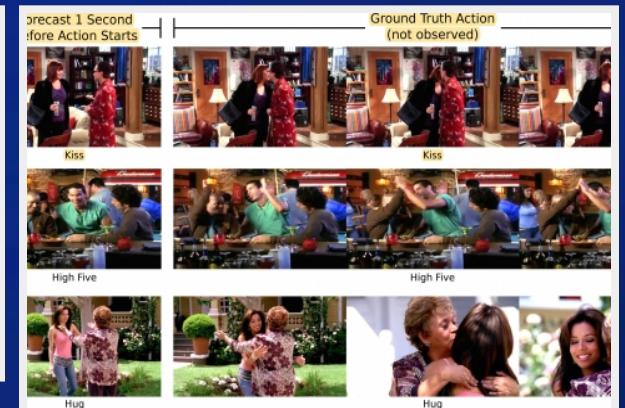
 **Yann LeCun**
@ylecun Following

And so it begins: convolutional nets built into ultrasound machine to help detect breast cancer.
fb.me/3Mdgib77W

RETWEETS 73 LIKES 76

6:04 PM - 22 Apr 2016

...



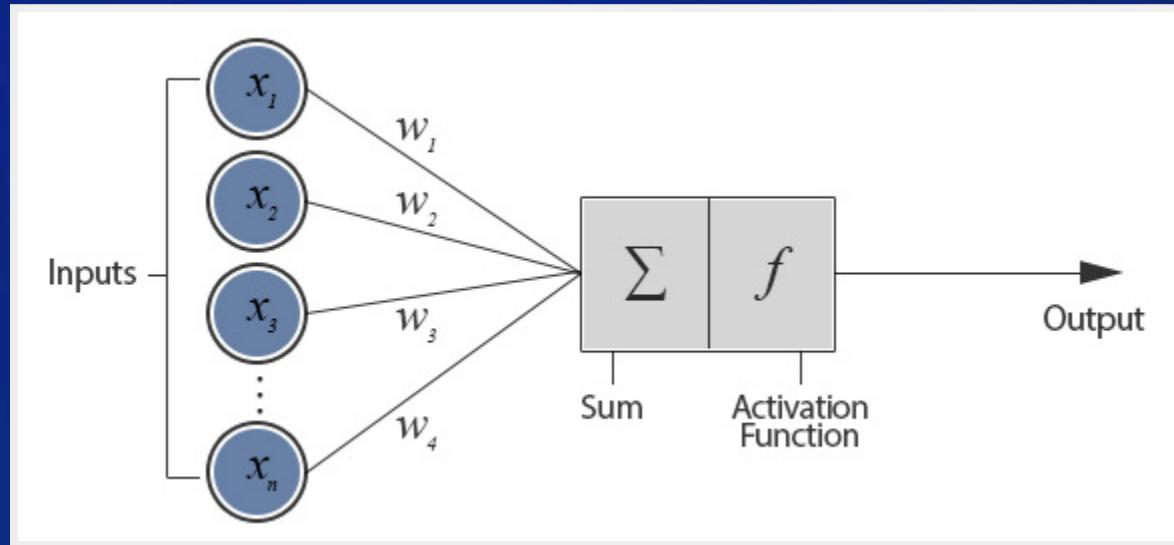
Why JavaScript for Machine Learning?

- Python and R are predominant
 - Have a large and mature set of libs
 - Are reasonably fast
 - Using binding to C/C++ or Fortran
- JavaScript has benefits, though
 - might be the language you are most comfortable with
 - might be the only language around (because all you have is a browser)
 - zero installation, easy to get started
 - combination with interactive visualizations

Visualizing the Basics

The perceptron - where it all begins

- mathematical model of a biological neuron
- creates a single output based on sum of many weighted inputs
- uses a **sigmoid** activation function to create boolean like output



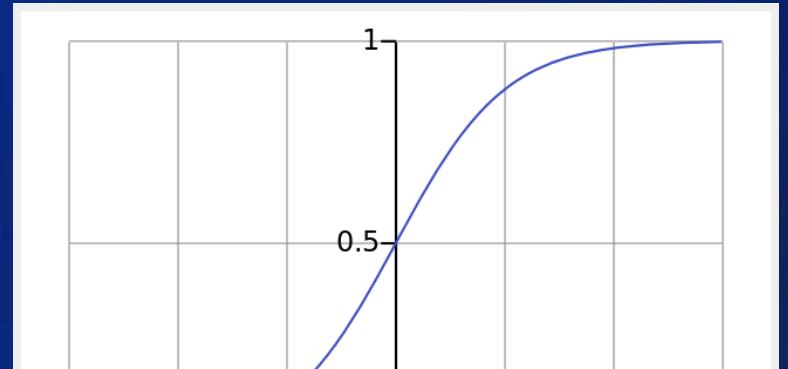
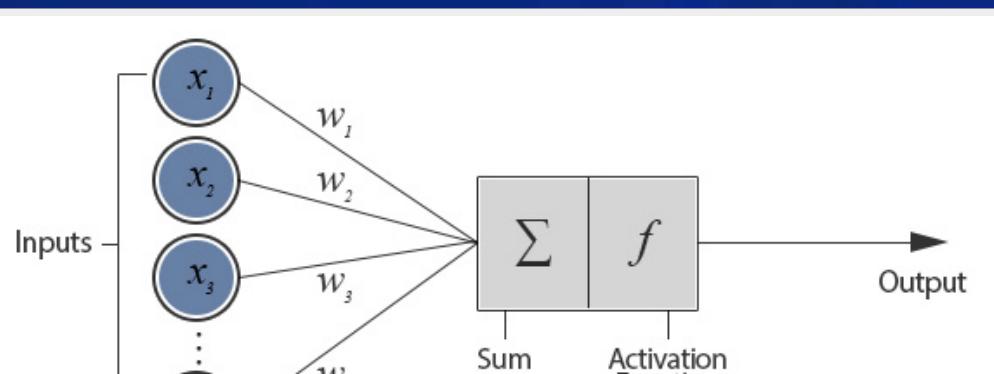
<http://www.theprojectspot.com/tutorial-post/introduction-to-artificial-neural-networks-part-1/7>

Implementing it in pure JavaScript

```
// Initial weights
let w0 = 3, w1 = -4, w2 = 2;

function perceptron(x1, x2) {
    const sum = w0 + w1 * x1 + w2 * x2;
    return sigmoid(sum);
}
```

```
function sigmoid(z) {
    return 1 / (1 + Math.exp(z * -1));
}
```

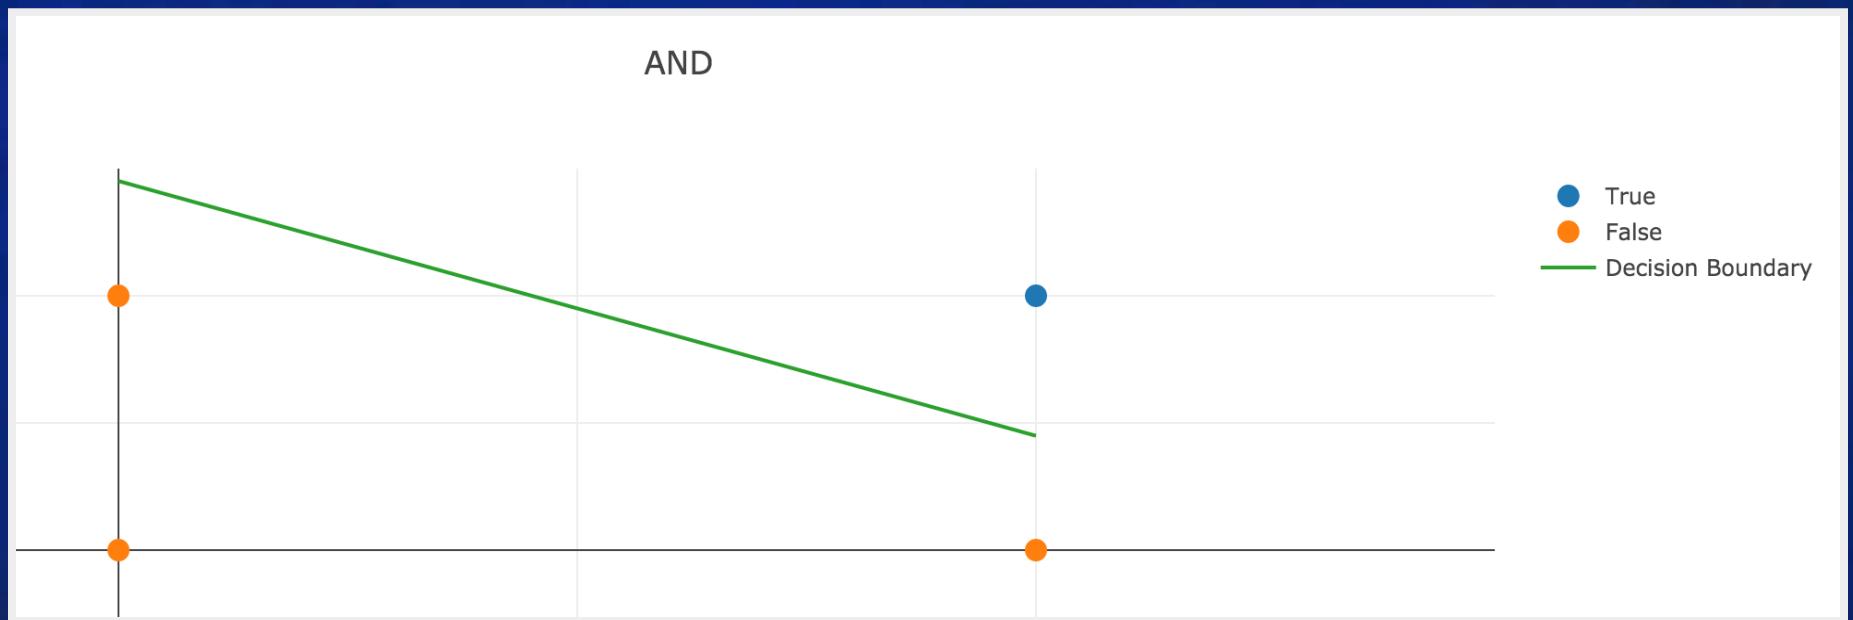


Right, that is surprisingly simple

Code for training is a little bit more complex (but only a little bit)

Visualizing what a neuron can do

- output separates plane into two regions using a line
- such regions are called *linearly separable*
- can emulate most logic functions (NOT, AND, OR, NAND)
- can be trained by adjusting weights of inputs based on error in output

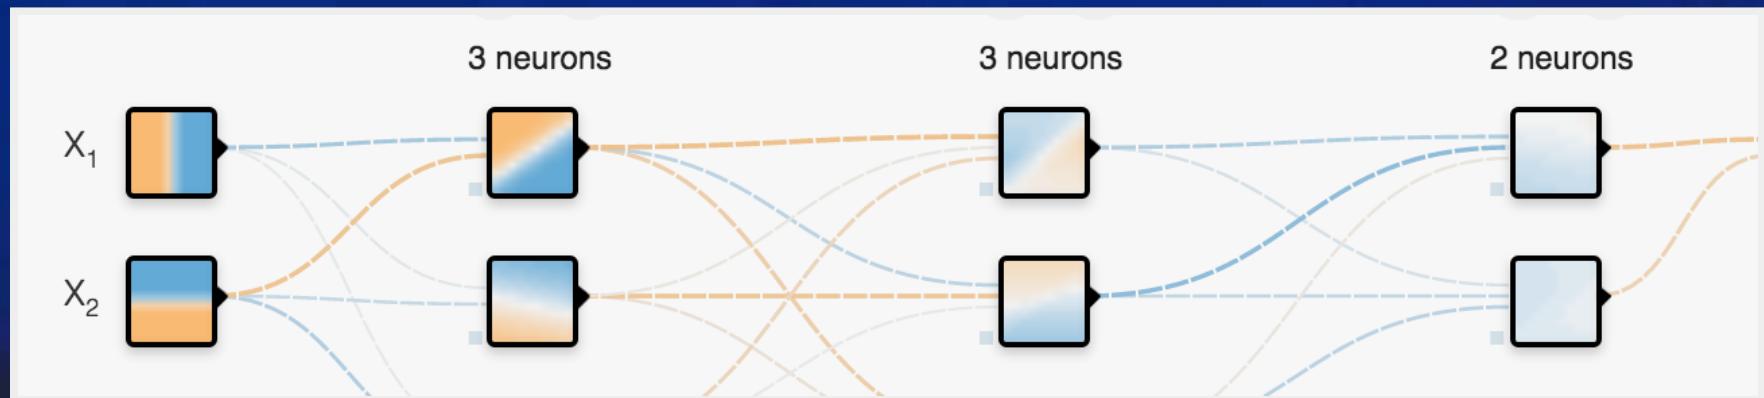


perceptron training visualization (initial version provided as a courtesy of [Jed Borovik](#))

**A single neuron is not very powerful
but becomes much more powerful when organized in layers**

Feedforward Neural Networks

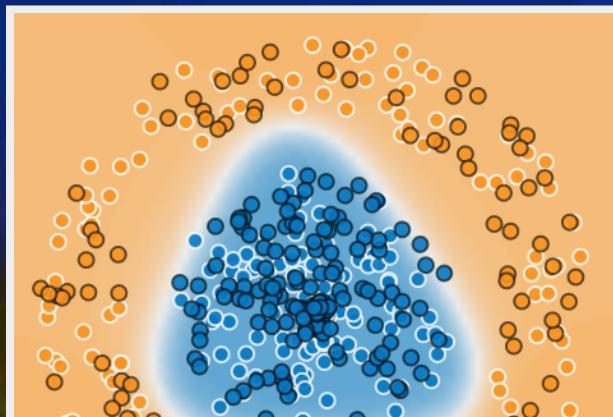
- neurons can be arranged in layers
- data flows forward in one direction, there are no cycles
- first layer takes inputs
- output layer (typically softmax) classifies summed up output
- middle layer(s) called hidden layer(s)
- more than one hidden layer is possible (deep neural network)
- each neuron in each hidden layer gets exactly the same input
- no connection between neurons in same layer
- each neuron in one layer feeds all neurons in the next layer



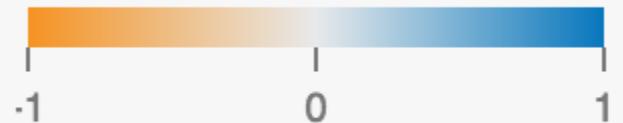
A Classification example

Using the Tensorflow Playground

- result of the neural network is a classification
- each spot in a plane gets a prediction
- either blue or orange
- prediction can be discrete or continuous
- network is trained to give a good prediction

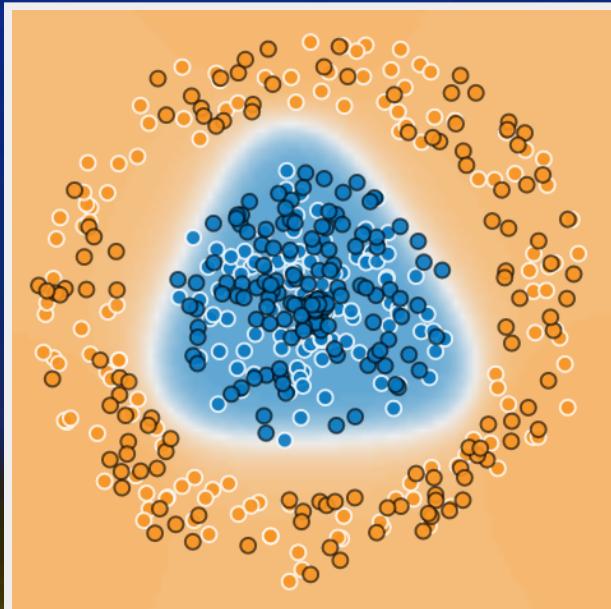


Colors shows
data, neuron and
weight values.

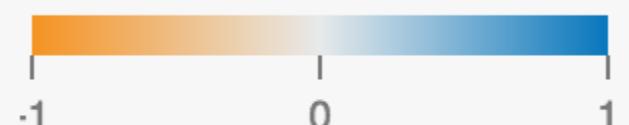


Training and Loss

- data is separated into training and test data
- in image below: training: white border, test: black border
- training data is used to train neural network
- both test and training data are used to check accuracy of prediction
- loss is to be minimized
- overfitting: training loss low, test loss much higher (to be avoided)



Colors shows
data, neuron and
weight values.



Tensorflow Internals

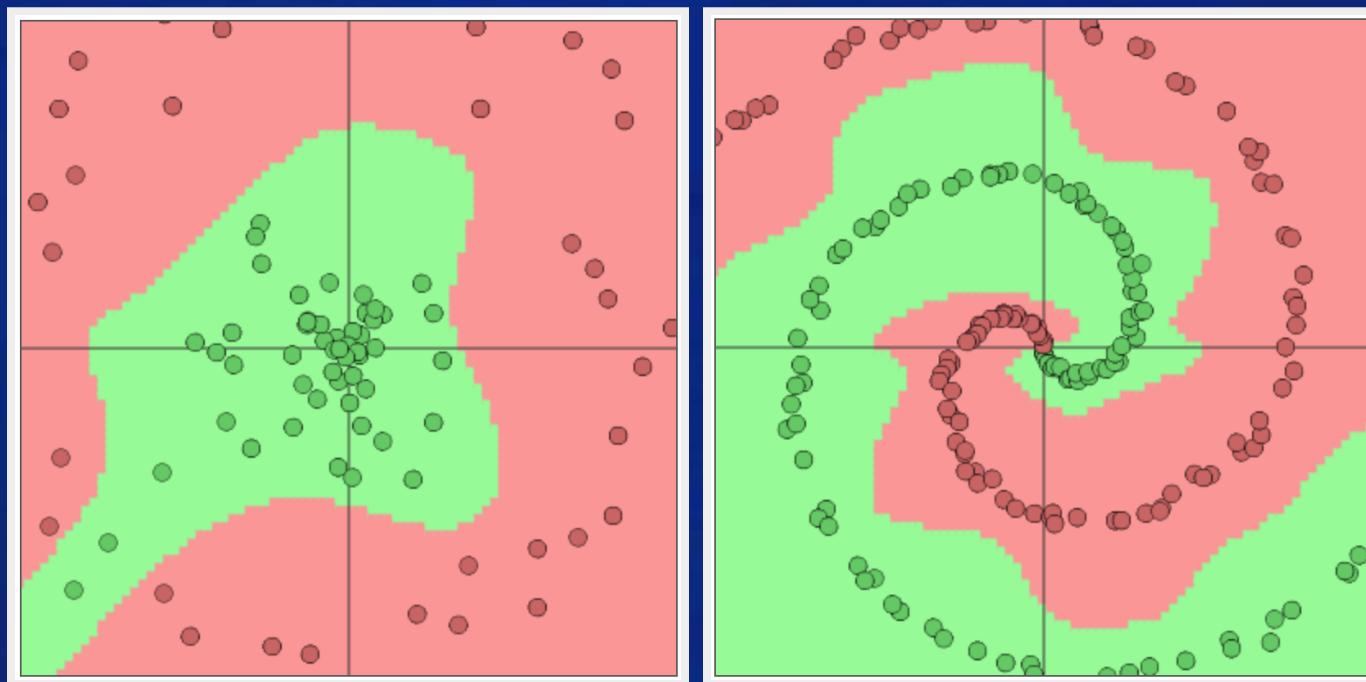
- Written in [TypeScript](#)
- Visualizations done with [D3.js](#)
- Everything implemented from scratch
- No performance optimizations
- Easy to understand
- NN simulation fairly similar to code from scratch as shown before
- More general for as many dimensions as you want

Now for the JavaScript

ConvNetJS

<http://cs.stanford.edu/people/karpathy/convnetjs/>

Code can be changed directly in the browser



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>

Code for classifier

That's what you can change directly in the browser

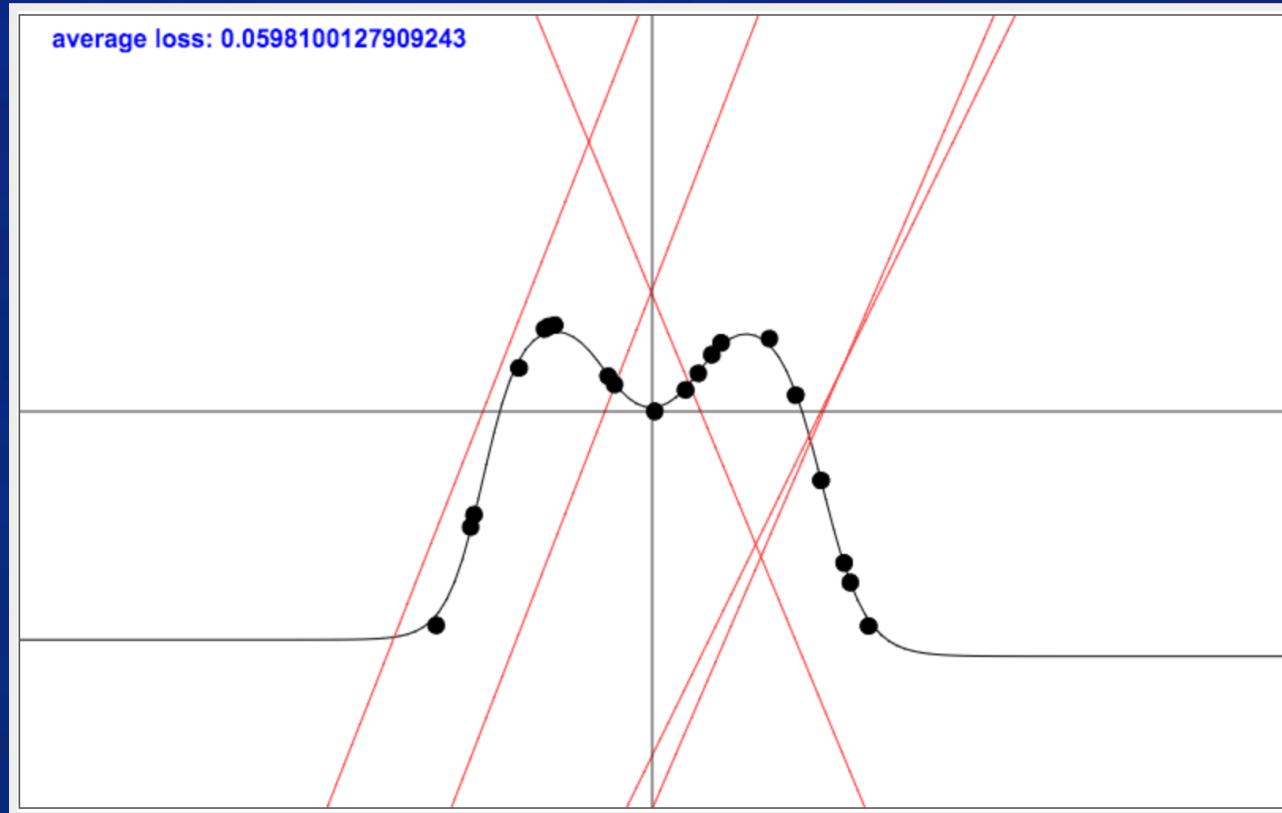
```
layer_defs = [
    {type:'input', out_sx:1, out_sy:1, out_depth:2},
    {type:'fc', num_neurons:6, activation:'tanh'},
    {type:'fc', num_neurons:2, activation: 'tanh'},
    {type:'softmax', num_classes:2}
];

net = new convnetjs.Net();
net.makeLayers(layer_defs);

trainer = new convnetjs.Trainer(net);
```

Regression

Does not classify, but tries to find a continuous function that goes through all data points



<http://cs.stanford.edu/people/karpathy/convnetjs/demo/regression.html>

Regression Example - Code

Quick Quiz: How many neurons in hidden layer?

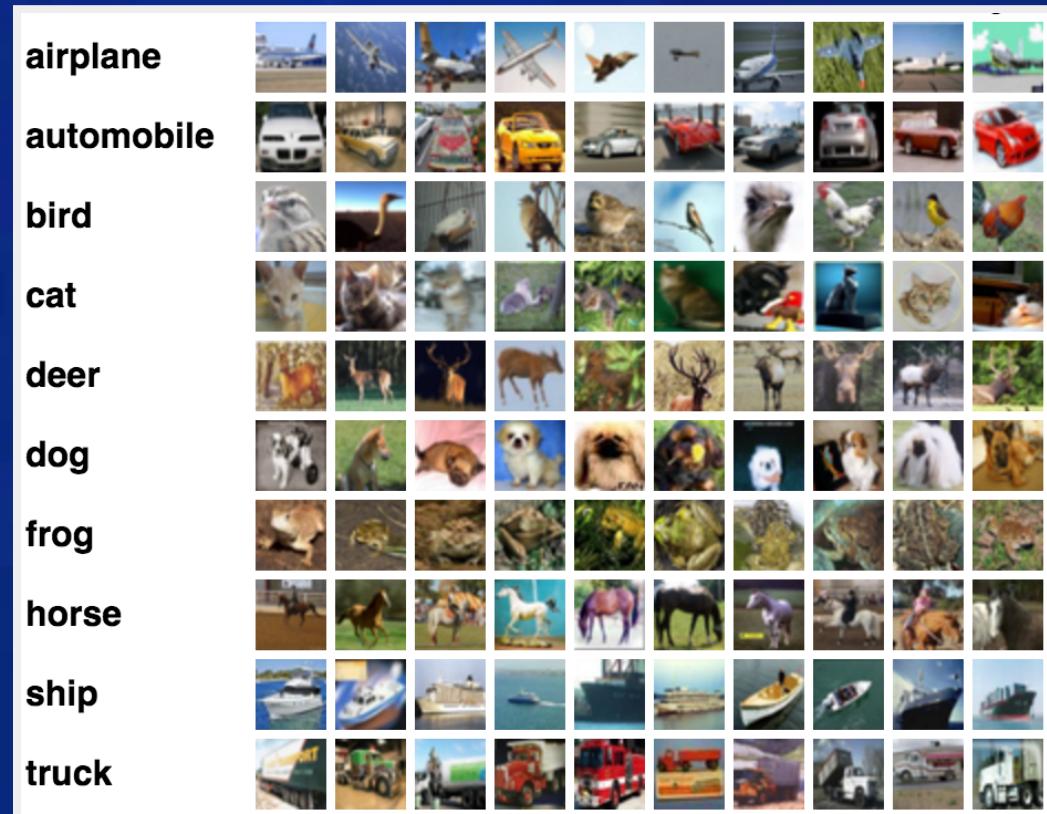
```
layer_defs = [  
    {type:'input', out_sx:1, out_sy:1, out_depth:1},  
    {type:'fc', num_neurons:5, activation:'sigmoid'},  
    {type:'regression', num_neurons:1}];
```

CIFAR-10 with Convolutional Deep Neural Networks

uses convolutional hidden layers for filtering

The CIFAR-10 dataset

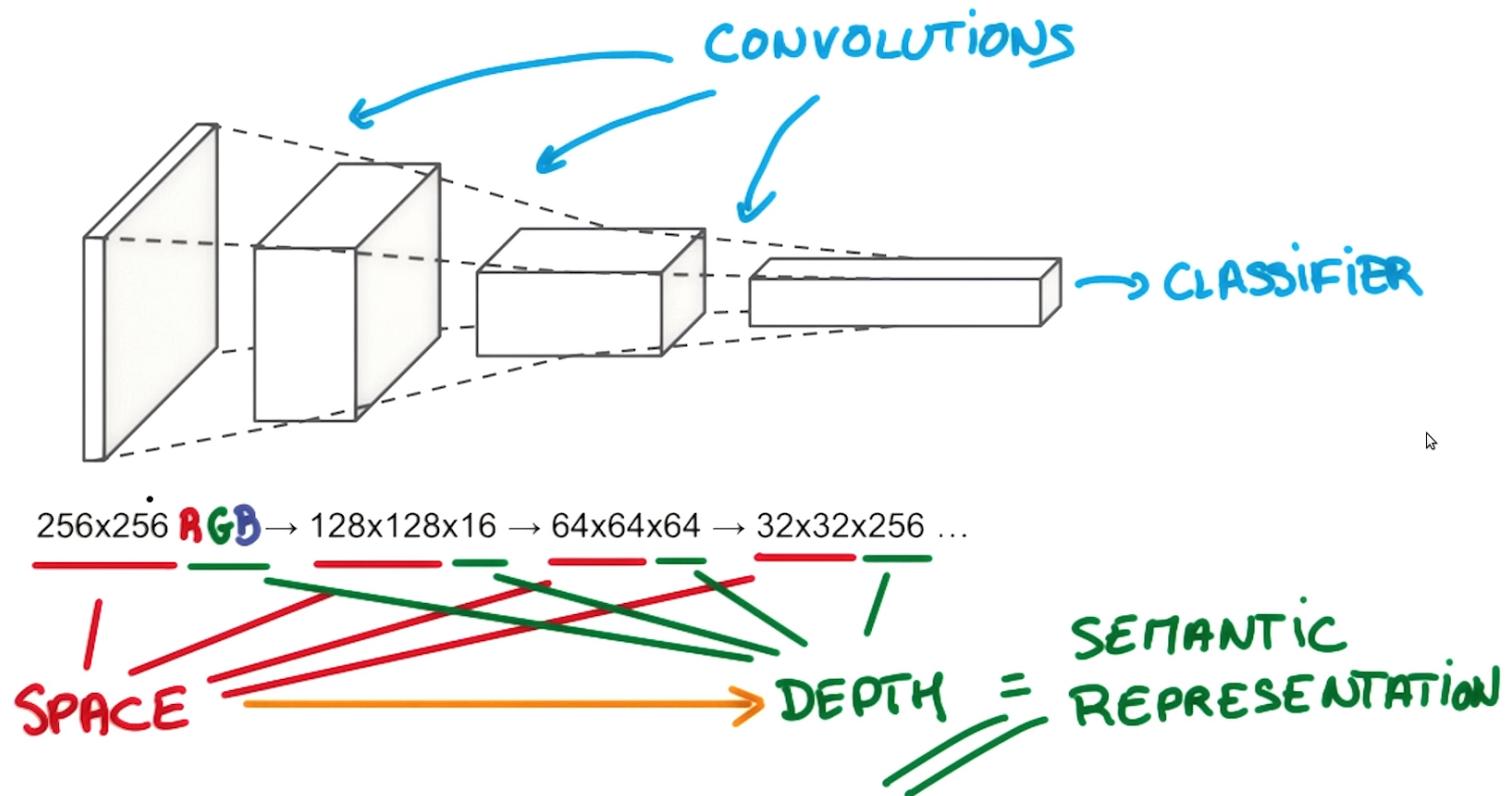
Tiny images in 10 classes, 6000 per class in training set



<http://www.cs.toronto.edu/~kriz/cifar.html>

The Convolutional Pyramid

CONVOLUTIONAL PYRAMID



CIFAR-10 Example - Code

```
layer_defs = [  
    {type:'input', out_sx:32, out_sy:32, out_depth:3},  
    {type:'conv', sx:5, filters:16, stride:1, pad:2, activation:'relu'},  
    {type:'pool', sx:2, stride:2}),  
    {type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'}.  
    {type:'pool', sx:2, stride:2},  
    {type:'conv', sx:5, filters:20, stride:1, pad:2, activation:'relu'},  
    {type:'pool', sx:2, stride:2},  
    {type:'softmax', num_classes:10}];
```

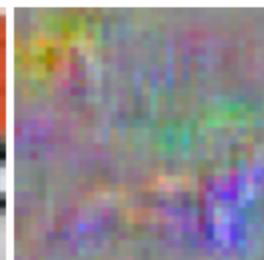
Input: 32x32 RGB images

input (32x32x3)

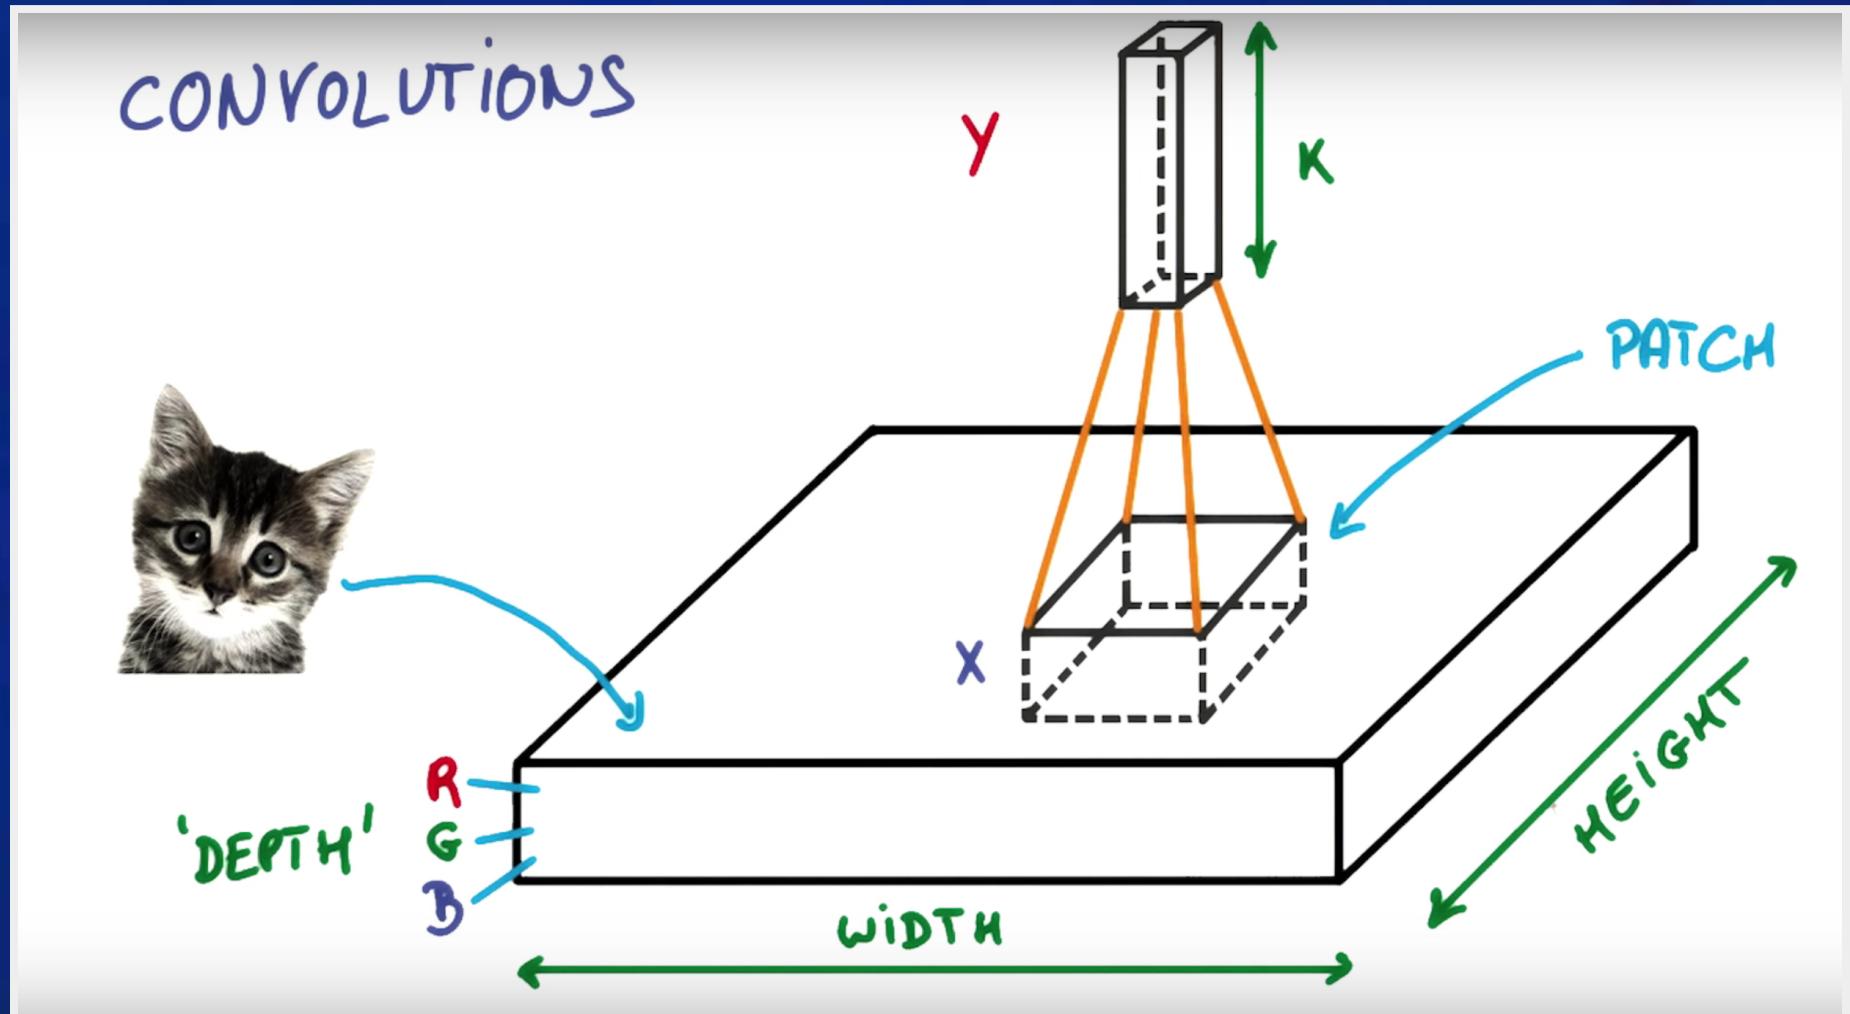
max activation: 0.5, min: -0.5

max gradient: 0.10624, min: -0.12467

Activations:



Convolutions



Convolution (Filtering) step #1: 16 32x32 images of filtered data

```
// 16 5x5 filters will be convolved with the input  
// output again 32x32 as stride (step width) is 1  
// 2 pixels on each edge will be padded with 0 matching 5x5 pixel filter  
{type:'conv', sx:5, filters:16, stride:1, pad:2, activation:'relu'}
```

conv (32x32x16)

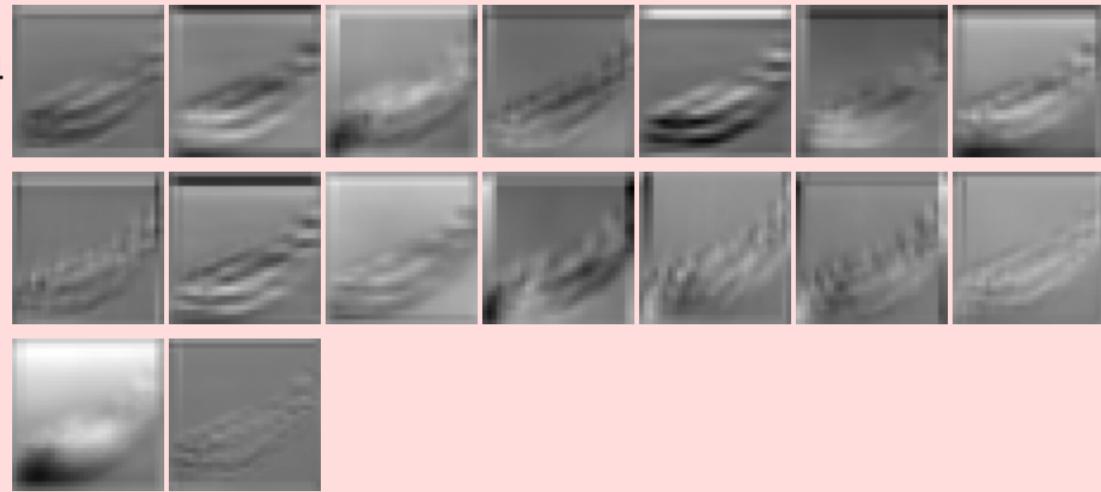
filter size 5x5x3, stride 1

max activation: 1.33221, min: -1.66454

max gradient: 0.01757, min: -0.01873

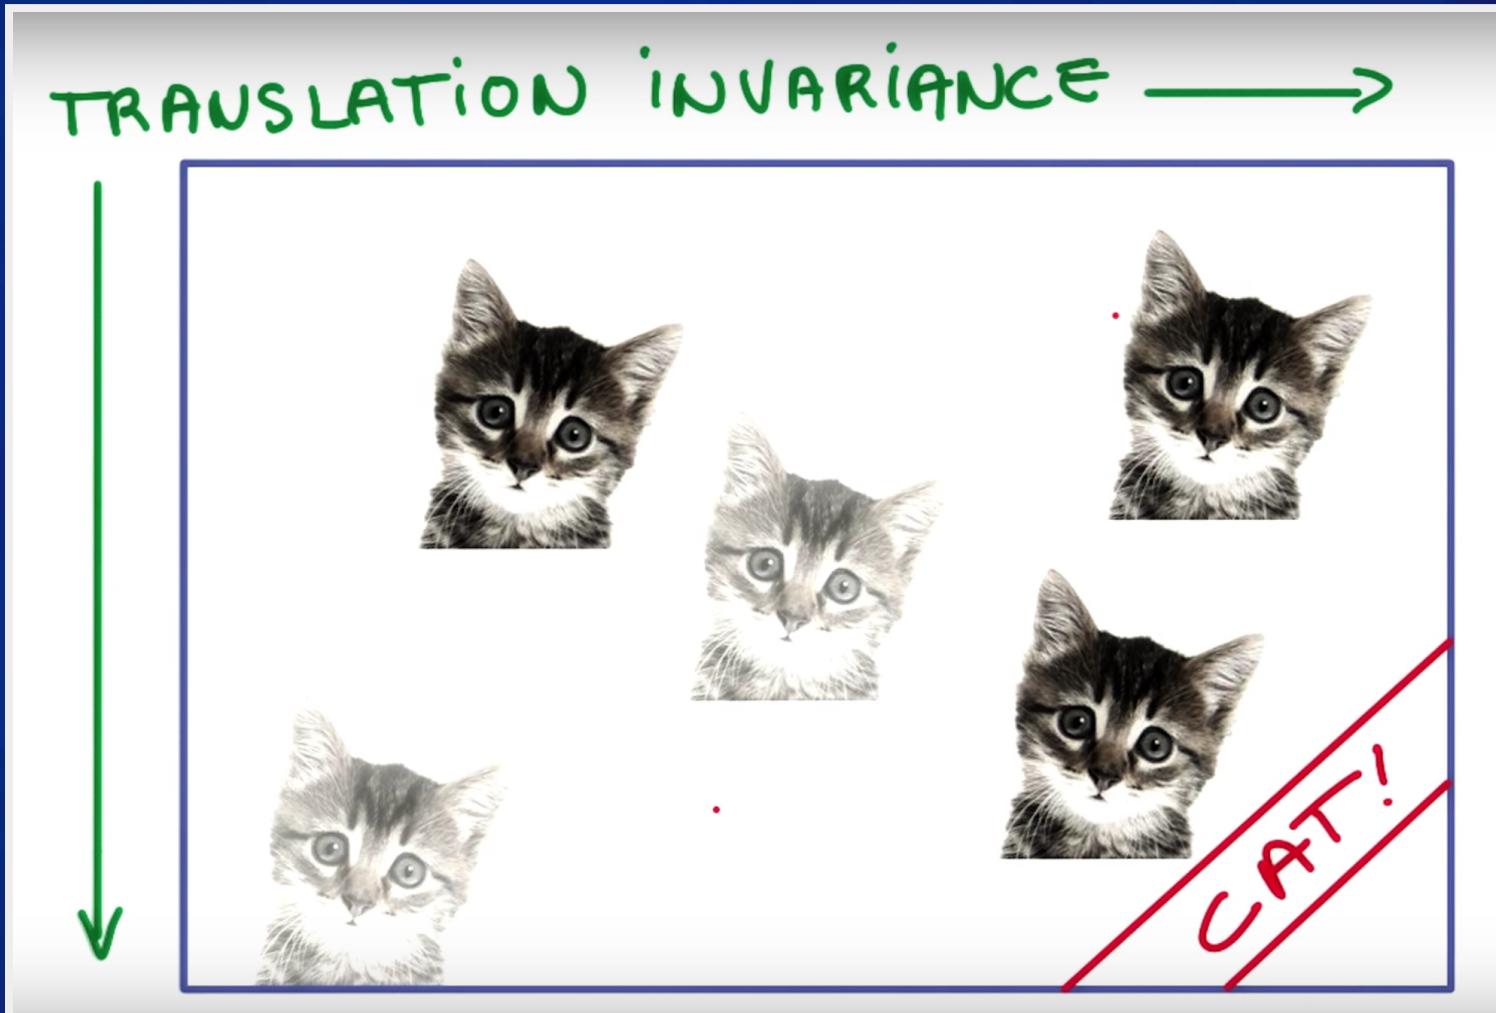
parameters: $16 \times 5 \times 5 \times 3 + 16 = 1216$

Activations:



convolutional layers remove noise, add semantics

Translation Invariance



Pooling

```
// perform max pooling in 2x2 non-overlapping neighborhoods  
// output will be 16x16 as stride (step width) is 2  
{type:'pool', sx:2, stride:2}),
```

pool (16x16x16)
pooling size 2x2, stride 2
max activation: 1.45481, min: 0
max gradient: 0.0798, min: -0.11926

Activations:



pooling layers in between provide translation invariance

Softmax

```
{type:'softmax', num_classes:10}
```

softmax (1x1x10)

max activation: 0.60745, min: 0.00011

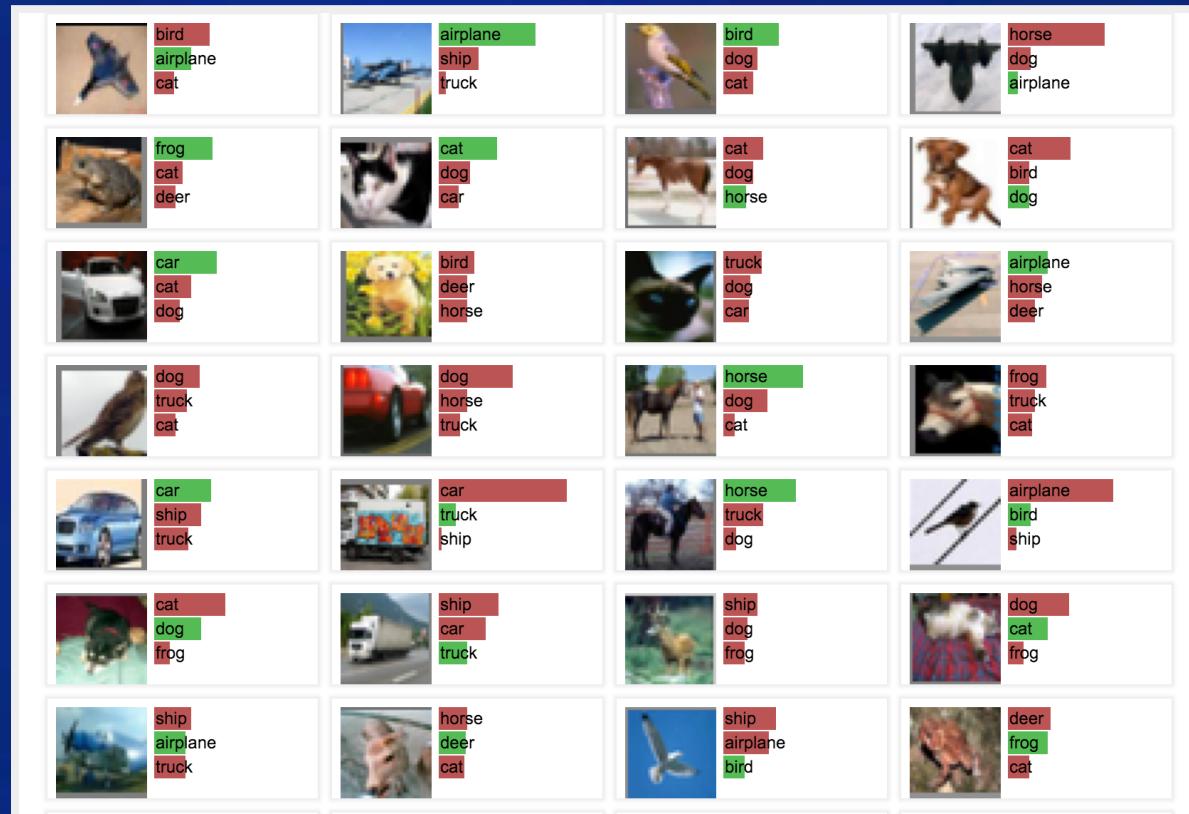
max gradient: 0, min: 0

Activations:



assigns a probability to each of the 10 categories

Sample classification results

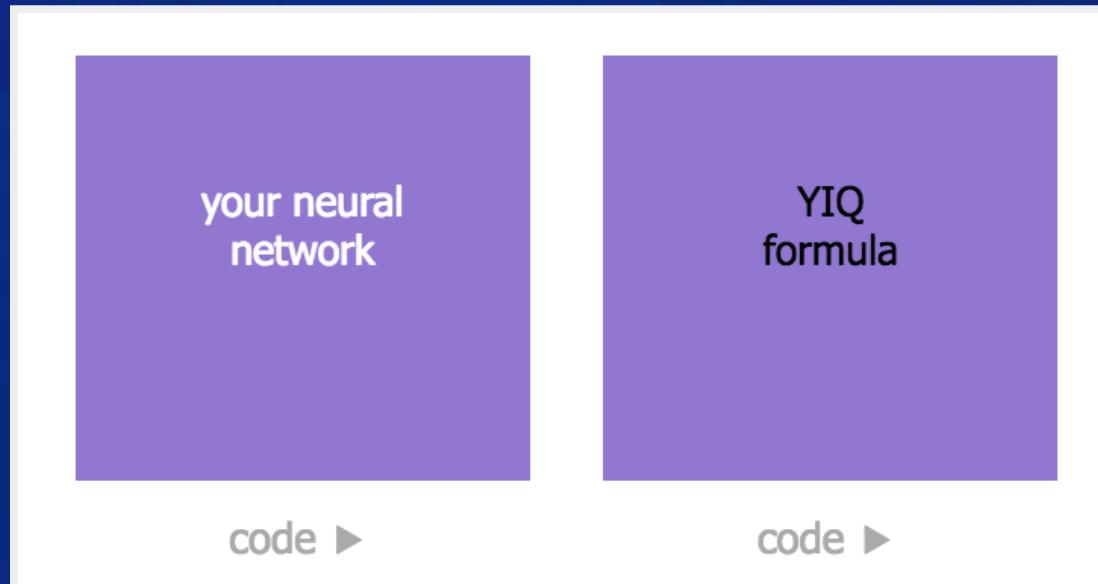


<http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

Other JavaScript libraries that run in the Browser

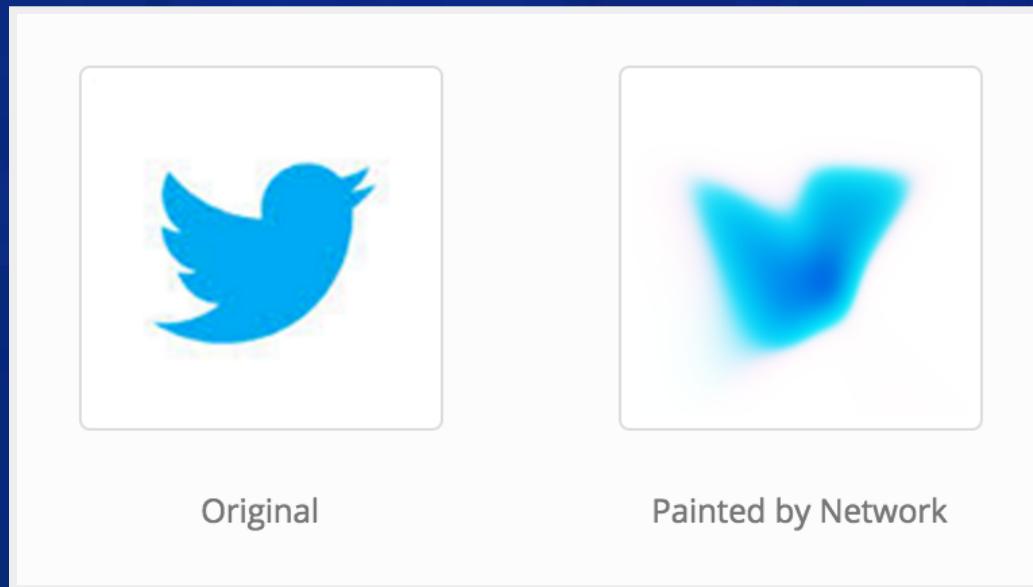
Brain.js

- very simple lib
- used to be unmaintained, but taken over by new maintainers



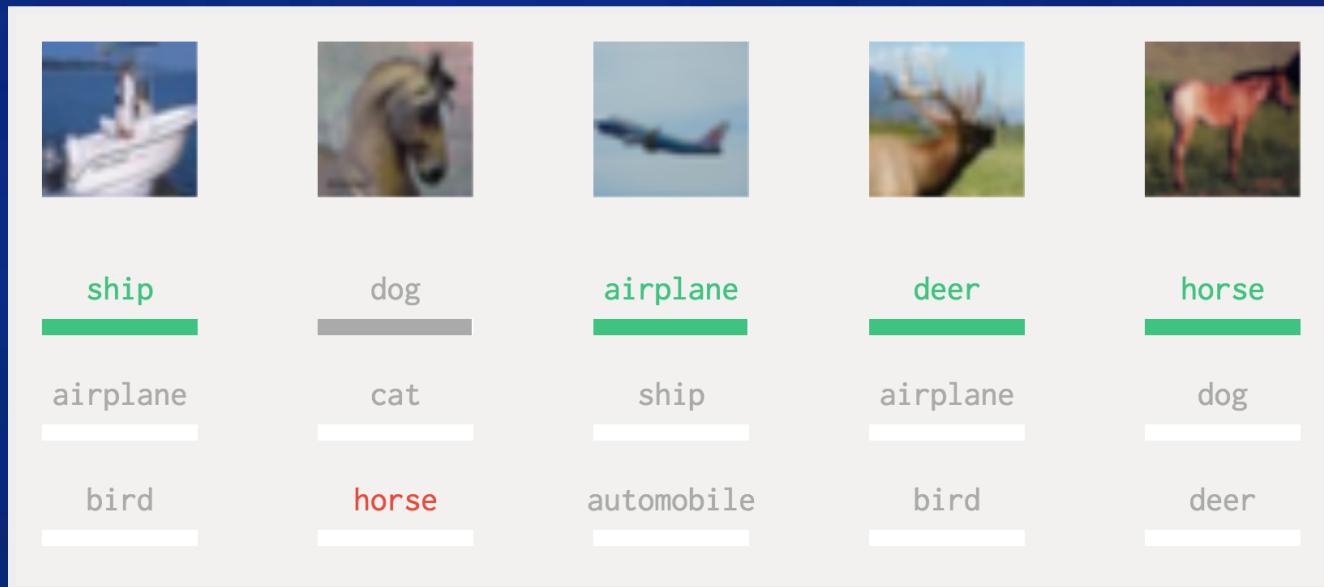
Synaptic.js

- small general lib for Neural Network
- nice and simple API
- good documentation
- actively maintained



neocortex.js

- Runs pre-trained deep neural networks (Keras models)
- Some really nice examples including Python Notebooks
- Clear code written in ES6



Wrap-Up

- Deep Neural Networks are a game changer in Machine Learning and AI
- enabled by large Data Sets, Computing Power, and GPUs
- the browser makes Deep Learning even more accessible
- all the fancy deep learning stuff works in the browser
- especially good for learning and education
- direct visualization and interactivity
- Tensorflow Playground makes this especially easy to try out
- ConvnetJS for experiments with JavaScript in the browser

Thank you!

Questions / Discussion

Oliver Zeigermann / @DJCordhose

<http://bit.ly/ml-data2day>